

The seL4 microkernel

Introduction to seL4 Microkernel

▶ What is seL4?

- ▶ A **formally verified**, high-performance operating system kernel.
- ▶ Designed for **security- and safety-critical systems**, embedded, and cyber-physical applications.

▶ Core Features

- ▶ Minimal **microkernel architecture**.
- ▶ Strong **isolation** and **fine-grained access control** through capabilities.
- ▶ Robust support for **real-time** and **mixed-criticality systems**.

▶ Key Features

- ▶ Minimal **Trusted Computing Base (TCB)** for reduced attack surface.
- ▶ Strong **isolation** through capability-based access control.
- ▶ **Formally verified** with machine-checked proofs
 - ▶ Ensures **functional correctness** and **security enforcement**.
 - ▶ Guarantees **confidentiality, integrity, and availability**.
- ▶ Verification spans from **abstract model** to **binary code**.

Introduction to seL4 Microkernel (cont'd)

- ▶ **Dual Role**

- ▶ Operating system **microkernel**.
- ▶ **Hypervisor** for secure virtual machines.

- ▶ **Real-Time and Mixed-Criticality Support**

- ▶ Complete analysis of **worst-case execution time**.
- ▶ Suitable for **hard real-time systems**.

Applications and Benefits of seL4

▶ Applications

- ▶ Autonomous vehicles, defense systems, and **embedded devices**.
- ▶ IoT devices requiring security and safety.
- ▶ Cyber retrofitting of legacy systems.

▶ Key Advantages

- ▶ Industry benchmark for **reliability** and **performance**.
- ▶ Enables integration of existing software into **secure environments**.
- ▶ Combines **robust security** with **practical adaptability**.

seL4 is a Microkernel, not an OS

- ▶ **Fundamental Difference** Unlike monolithic kernels (e.g., Linux), seL4
 - ▶ Minimizes code in **privileged mode**.
 - ▶ Reduces the **Trusted Computing Base (TCB)** and **attack surface**.

Monolithic v. microkernel

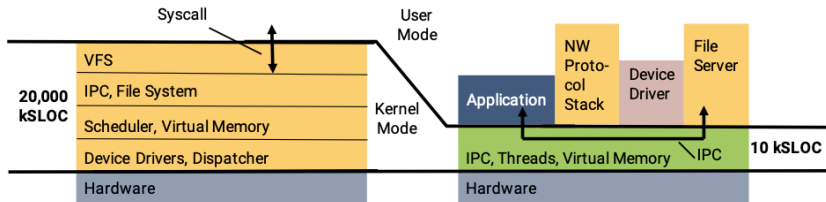


Figure 1: Operating system structure: Monolithic v. microkernel

Monolithic vs. Microkernel

▶ **Monolithic Kernels**

- ▶ Integrate all essential services (e.g., drivers, file systems) into the kernel.
- ▶ Tens of millions of lines of code = higher vulnerability.

▶ **Microkernels**

- ▶ Provide only minimal functionality for managing hardware and isolating processes.
- ▶ Delegate OS services to **user-space programs**.

▶ **Benefits of seL4**

- ▶ Strong **isolation**.
- ▶ Fine-grained access control using **capabilities**.
- ▶ Lightweight, low-level API for high efficiency.
- ▶ Modular design ensures **fault isolation**, **security**, and **resilience**.

Hypervisor Capabilities of seL4

► Virtual Machine Support

- Secure execution of full OSES (e.g., Linux) alongside **native applications**.
- Enables **seamless integration** of native and virtualized components.

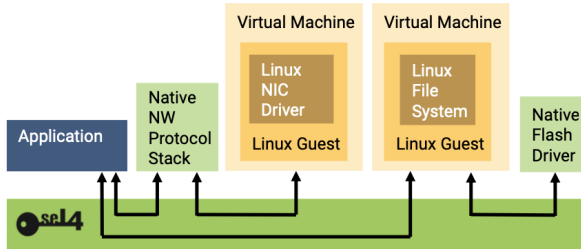


Figure 2: Virtualization to integrate native services with Linux

▶ Key Features

▶ Fine-grained capability-based access control

- ▶ Ensures strict isolation between VMs, native services, and applications.
- ▶ Prevents faults or compromises in one domain from affecting others.
- ▶ Secure interaction via **well-defined communication channels**.

▶ Incremental Modernization

- ▶ Run **legacy systems** in virtual machines.
- ▶ Operate new components **natively** for enhanced security and performance.
- ▶ Achieves modernization without a full system overhaul.

seL4 for Real-Time Systems

- ▶ **Precise Timing Guarantees**
 - ▶ Priority-based scheduling for **predictable execution**.
 - ▶ Developers can control thread priorities to meet **strict deadlines**.
- ▶ **Key Features for Real-Time**
 - ▶ **Bounded worst-case execution times** for all kernel operations.
 - ▶ **Minimal and predictable interrupt latencies**, even under heavy workloads.
- ▶ **Mixed-Criticality Systems (MCS)**
 - ▶ Secure coexistence of components with varying safety and timing needs.
 - ▶ Strong isolation prevents interference between components.
- ▶ **Capability-Based Resource Management**
 - ▶ Precise, secure allocation of **time resources** alongside memory and I/O.
- ▶ **Ideal Applications**
 - ▶ Avionics, autonomous vehicles, and time-critical **embedded systems**.

Verification of seL4

► **Formal Verification**

- First OS kernel with **machine-checked verification** of functional correctness.
- Ensures implementation aligns with high-level specifications.
- Guarantees absence of
 - Buffer overflows.
 - Null-pointer dereferences.
 - Code injection vulnerabilities.

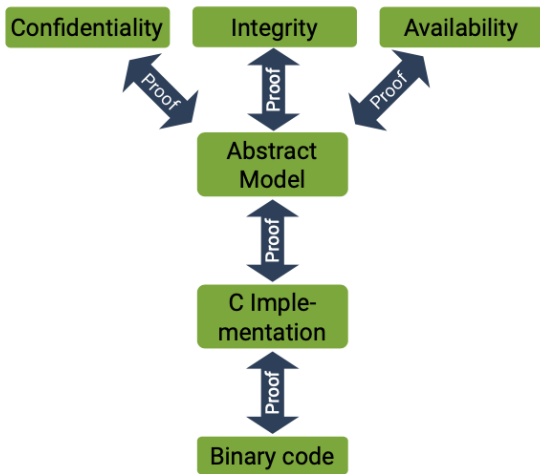


Figure 3: seL4 proof chain

▶ **Translation Validation**

- ▶ Extends verification to the compiled binary.
- ▶ Proves that the binary matches the verified source code, even with optimizations.
- ▶ Ensures **high assurance** in the deployed form.

▶ **Security Properties**

- ▶ Proven enforcement of **confidentiality, integrity, and availability**.
- ▶ Strict access controls and effective isolation of components.

▶ **Assumptions for Verification**

- ▶ Correctness of hardware.
- ▶ Accuracy of specifications.
- ▶ Reliability of the theorem prover.

▶ **Impact**

- ▶ Bridges formal reasoning and real-world execution.
- ▶ Sets a **new benchmark** for security and reliability in critical systems.

Functional Correctness of seL4

► Definition

- Rigorous proof that seL4's C implementation - Is free from defects. - Adheres to a formal specification expressed in **higher-order logic (HOL)**.
- Guarantees the kernel behaves strictly as defined by its **abstract model**.

► Verification Process

- Uses **Isabelle/HOL** theorem prover.
- Translates C code into mathematical logic for formal verification.
- Restricts C usage to a well-defined subset with unambiguous semantics.
- Ensures the implementation remains **provably correct**.

Translation Validation in seL4

Why Translation Validation?

- ▶ Bug-free C implementation \leftrightarrow Guaranteed reliability
 - ▶ **Compilers** are complex systems that may
 - ▶ Introduce defects during code translation.
 - ▶ Contain bugs or malicious code (e.g., Trojan backdoors described by Ken Thompson).

What is Translation Validation?

- ▶ Process of verifying the compiled binary against the **formally verified C code**.
- ▶ Ensures the binary faithfully represents the verified source code.

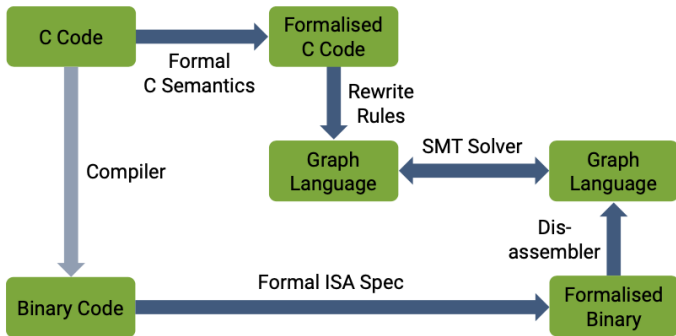


Figure 4: Translation validation proof chain

▶ **How it Works**

1. Formalization of the processor's ISA (instruction set architecture).
2. Disassembly of the binary and transformation into a graph-based intermediate representation.
3. Transformation of C code into the same intermediate representation.
4. Use of SMT solvers and rewrite rules to prove equivalence.

▶ **Key Outcome**

- ▶ Validates that the compiler's output matches the **abstract specification**.
- ▶ Ensures seL4's high-assurance guarantees extend to the executable binary.

▶ **Impact**

- ▶ Bridges the gap between the **formal model** and **real-world deployment**.
- ▶ Protects against risks introduced by compiler behavior.

Security Properties of seL4

- ▶ **Core Guarantees**

- ▶ **Confidentiality**

- ▶ Ensures no unauthorized entity can read or infer data.
 - ▶ Enforced through strict access control mechanisms.

- ▶ **Integrity**

- ▶ Prevents unauthorized modifications to data.

- ▶ **Availability**

- ▶ Protects against denial of authorized resource access.

- ▶ **Formal Proofs**

- ▶ Demonstrate the kernel's ability to secure critical systems.
 - ▶ Mitigate common attack vectors.

- ▶ **Limitations**

- ▶ Current model does **not yet cover timing-related security issues**.
 - ▶ Timing channels remain an area of active research.

- ▶ **Enhancements for Real-Time Systems**

- ▶ **Mixed-Criticality Systems (MCS)** model
 - ▶ Extends integrity and availability guarantees to include **timeliness**.
 - ▶ Ensures security in real-time environments.

Proof Assumptions in seL4 Verification

► **Explicit Assumptions in Formal Reasoning**

- Every assumption is **explicitly defined** and documented.
- Prevents risks of **overlooking or misinterpreting** critical dependencies.
- Enhances clarity and confidence in system correctness.

► **Key Assumptions**

1. **Hardware behaves as expected**

- Kernel guarantees depend on reliable hardware.
- Faulty or malicious components invalidate kernel assurances.

2. **Specification matches expectations**

- Formal specification must align with intended behavior.
- A gap may exist between **mathematical reasoning** and **real-world interpretation**.

3. **Theorem prover is correct**

- Tools like Isabelle/HOL have a **small, well-tested core**.
- The risk of bugs in the prover affecting proofs is extremely low.

CAmkES Component Framework

- ▶ **What is CAmkES?**

- ▶ A framework for designing systems on **seL4** as collections of **isolated components**.
- ▶ Components interact through **defined communication channels**.
- ▶ Formal **Architecture Description Language (ADL)** ensures secure system interactions.

Main Abstractions in CAmkES

1. Components

- ▶ Represented as **square boxes**.
- ▶ Self-contained units of **code and data**.
- ▶ Encapsulated by **seL4**, functioning as independent programs.

2. Interfaces

- ▶ Define how components interact
 - ▶ **Importing** Invoke another component's interface.
 - ▶ **Exporting** Allow other components to invoke their interface.
 - ▶ Symmetric **shared-memory interfaces** for direct data sharing.

3. Connectors

- ▶ Link **importing** and **exporting interfaces** for communication.
- ▶ **One-to-one** by design; additional components enable **broadcast** or **multicast**.

Automated Translation and Setup

- ▶ **ADL** (architecture description language) CapDL** (capability distribution language)
 - ▶ ADL specifies architecture; CapDL defines **seL4 objects** and **access rights**.
 - ▶ Ensures faithful implementation of the described architecture.
- ▶ **Generated Code**
 - ▶ **Startup code** Initializes seL4 objects and allocates capabilities.
 - ▶ **Glue code** Simplifies communication between components via function calls.

Key Benefits

- ▶ Simplifies **design**, **verification**, and **implementation** of secure systems.
- ▶ Enables **sandboxed components** with precise communication channels.
- ▶ Maintains **security and reliability** through formal specifications and automated tools.

Capabilities in seL4

► What are Capabilities?

- Object references similar to **pointers** but include **access rights**.
- Immutable and uniquely reference specific objects.
- Encapsulate the rights needed to operate on objects.



Figure 5: A capability is a key that conveys specific rights to a particular object

Example

- ▶ An operation may be to call a function in a component.
- ▶ The object reference embedded in the capability then points to an interface to that object
- ▶ It conveys the right to invoke that function, a particular method on the component object.
- ▶ The capability may or may not at the same time convey the right to pass another capability along as a function argument
 - ▶ (delegating to the component the right to use the object referenced by the capability argument).

Key Features of Capabilities

1. Fine-Grained Access Control

- ▶ Invoking a capability is the only way to operate on system objects.
- ▶ Ensures strict adherence to the **principle of least privilege**.

2. Delegation

- ▶ Capabilities can be passed to delegate access securely.

3. Kernel Protection

- ▶ Unlike traditional ACLs, capabilities are
 - ▶ Managed by the **kernel**.
 - ▶ Immune to vulnerabilities like the **confused deputy problem**.

Types of Objects in seL4 referenced capabilities

1. **Endpoints** are used to perform protected function calls;
2. **Reply Objects** represent a return path from a protected procedure call;
3. **Address Spaces** provide the sandboxes around components (thin wrappers abstracting hardware page tables);
4. **Cnodes** store capabilities representing a component's access rights;
5. **Thread Control Blocks** represent threads of execution;
6. **Scheduling Contexts** represent the right to access a certain fraction of execution time on a core;
7. **Notifications** are synchronisation objects (similar to semaphores);
8. **Frames** represent physical memory that can be mapped into address spaces;
9. **Interrupt objects** provide access to interrupt handling; and
10. **Untyped** unused (free) physical memory that can be converted ("retyped") into any of the other types.

Benefits of Capabilities

- ▶ **Security**

- ▶ Restrict access to the **minimum rights** required for tasks.
- ▶ Avoid common vulnerabilities of traditional systems.

- ▶ **Comprehensive Control**

- ▶ Offers precise, object-oriented access control.

- ▶ **Reliability**

- ▶ Robust enforcement of system policies ensures high assurance.

Fine-Grained Access Control in seL4

- ▶ **Capabilities vs. ACLs** (access control list)
 - ▶ **Capabilities**
 - ▶ Object-oriented access control.
 - ▶ Aligns with the **Principle of Least Privilege (POLA)**.
 - ▶ Grants access only to explicitly authorized resources.
 - ▶ **ACLs** (Access-Control Lists)
 - ▶ Subject-oriented (based on user or group IDs).
 - ▶ Coarse-grained permissions limit precise security enforcement.

Access control in Linux

- ▶ File has an a set of access-mode bits.
 - ▶ Some of those bits determine what operations its owner can perform on the file
 - ▶ Others represent operations by each member of the file's "group"
 - ▶ A third gives rights to everyone else

```
-rw-r--r--@    1 dgcole  staff    181 Nov 12 14:48 .zshrc
drwx-----@    4 dgcole  staff    128 Aug 27 13:26 Applications/
drwxr-xr-x@   31 dgcole  staff    992 Jan 22 17:57 Desktop/
```

- ▶ This is a **subject-oriented** theme
 - ▶ a property of the subject/process attempting access
 - ▶ all subjects with the same value of the property (ID or group) have the same rights

Capability control

- ▶ Capabilities provide object-oriented access control
- ▶ The kernel will allow an operation if and only if the subject requesting the operation presents a capability that allows it to perform the operation
- ▶ E.g., an untrusted app can only access files for which it has been given a capability
- ▶ Alice invokes a program
 - ▶ she handing it a capability to the one files the program is allowed to read,
 - ▶ plus a capability to a file where the program can write its output,
 - ▶ the program is unable to access anything else — proper least privilege.

Confinement of Untrusted Programs

▶ **Traditional Systems (e.g., Linux)**

- ▶ Restricting access requires cumbersome workarounds like
 - ▶ chroot jails
 - ▶ Containers

▶ **Capabilities in seL4**

- ▶ Precisely grant access to specific resources
 - ▶ Example: Reading/writing specific files.
- ▶ Ensure untrusted programs cannot interact with unauthorized resources.

Key Benefits of Capabilities

1. **Precision**

- ▶ Allows application-specific access control.

2. **True Least Privilege**

- ▶ Confines programs to minimal required permissions.

3. **Simplicity and Security**

- ▶ Avoids the complexity and vulnerabilities of ACL-based systems.

Solutions for Delegation and Interposition

Interposition

- ▶ Capabilities enable transparent mediation
 - ▶ Capabilities are opaque references.
 - ▶ Example: A capability given to a user points to a security monitor instead of the resource.
- ▶ Applications of interposition
 - ▶ Enforcing security policies.
 - ▶ Packet filtering.
 - ▶ Debugging and resource virtualization.

Interpose access

A consequence of the fact that they are opaque object references. I

- ▶ Alice is given a capability to an object, she has no way of knowing what that object really is, all she can do is invoke methods on the object.
- ▶ The system designer may pretend that the capability given to Alice refers to a file, when in fact it refers to a communication channel to a security monitor, which in turns holds the actual file capability.
- ▶ The monitor can examine Alice's requested operations
 - ▶ If valid,
 - ▶ then performs them on the file on her behalf
 - ▶ else ignoring invalid ones
- ▶ The monitor effectively virtualises the file.

Delegation with Capabilities

▶ **Efficient Privilege Delegation**

- ▶ Users can “**mint**” new capabilities with specific permissions (e.g., **read-only access**).
- ▶ Delegate capabilities securely to others.
- ▶ **Revocation**
 - ▶ Capabilities can be revoked at any time, enhancing control.

▶ **Autonomous Resource Management**

- ▶ Subsystems can independently manage their resources
 - ▶ Maintain isolation and security.
 - ▶ Avoid reliance on centralized control.

Delegation

Capabilities support safe and efficient delegation of privilege.

- ▶ If Alice wants to give Bob access to one of her objects,
 - ▶ then she can create (“mint” in seL4 speak) a new capability to the object and hand it to Bob.
- ▶ Bob then can use that capability to operate on the object without referring back to Alice.
- ▶ The new capability can have diminished rights
- ▶ Alice can use this to give Bob only read-only access to the file
- ▶ Alice can revoke Bob’s access at any time by destroying the derived capability she handed to Bob.

Delegation is powerful and cannot easily and safely be done in ACL systems

A typical use case: setting up sub-systems that manage resources autonomously

- ▶ When the system starts up, the initial process holds authority to all resources in the system
- ▶ This initial resource manager can then partition the system, by creating new processes (secondary resource managers) and handing them privilege to disjoint subsets of the system resources
- ▶ The subsystems can then autonomously, without referring back to the original manager, control their subset of resources, while unable to interfere with each other
- ▶ Only if they want to change the original resource allocation do they need to involve the original manager.

Benefits of Capabilities for Delegation and Interposition

1. **Flexibility**

- ▶ Tailor capabilities for specific tasks or permissions.

2. **Transparency**

- ▶ Mediate access without revealing the resource.

3. **Control**

- ▶ Simplify revocation and enhance autonomous management.

4. **Enhanced Security**

- ▶ Achieve goals difficult with traditional access-control systems.

The Confused Deputy Problem

- ▶ **Definition**

- ▶ A security vulnerability where a program (**deputy**) is tricked into misusing its authority, leading to unintended actions.

Classic Example The Compiler Incident

A C compiler.

- ▶ It takes a C source file and produces an object-code output file, the file names are passed as arguments
- ▶ A user, Alice, must have execute permission on the compiler

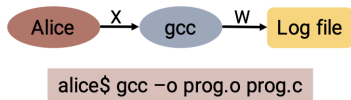


Figure 6: The compiler as a confused deputy

- ▶ Assume the compiler also creates an entry in a system-wide log file for auditing purposes
- ▶ The log file is not accessible to normal users, so the compiler must execute with elevated privilege in order to write to the log file

If Alice is malicious,

- ▶ then she can specify the password file as the output file
- ▶ The compiler will just open the output file (password file) and overwrite it with the compiled object code
- ▶ It doesn't take a lot of skill for Alice to write a program that compiles a password file will give her privileges she should not have

The fundamental problem: ACL-based systems use ambient authority for determining access rights.

- ▶ When the compiler opens its output to write, the OS looks at the compiler's subject ID to determine whether it has access to the object
- ▶ It is up to the compiler to determine whether the operation is valid or not, making the compiler part of the system's TCB
- ▶ It has to be fully trusted to do the right thing under all circumstances

Root Cause

The confusion arises due to ambient authority:

- ▶ The validity of an operation is determined by the security state of the agent (compiler),
 - ▶ a deputy operating on behalf of an original agent (Alice).
- ▶ For proper security, the access must be determined by Alice's security state.
 - ▶ This means that denomination (the reference to the file) and authority (the right to perform operations on the file) must be coupled
 - ▶ Principle called **no designation without authority**
- ▶ If that is the case, then the compiler invokes the designated object (output file) with the authority that comes with the designation (from Alice), and Alice can no longer confuse the deputy.

This is exactly what a capability system enforces.

► **Solution**

- Use capabilities that combine object designation with access rights.
- Ensures programs operate only within explicitly granted permissions, preventing such vulnerabilities.

Alice needs to hold three capabilities:

1. an execute capability on the compiler,
 2. a read capability on the input file, and
 3. a write capability on the output file.
- ▶ She invokes the compiler with the execute capability and passes the other two as arguments.
 - ▶ When the compiler then opens the output file, it does so with the capability provided by Alice, and there is no more confusion possible.
 - ▶ The compiler uses a separate capability, which it holds itself, for opening the log file, keeping the two files well separated.

It is impossible for Alice to trick the compiler into writing to a file she has no access to herself.

Benefits of Capabilities in Avoiding Confusion

1. **Elimination of Ambient Authority**

- ▶ Prevents unauthorized operations.

2. **Secure Delegation**

- ▶ The user explicitly controls access rights.

3. **True Principle of Least Privilege**

- ▶ Programs act only on resources for which they have explicit capabilities.

4. **Essential for Secure Environments**

- ▶ Ensures reliable and unambiguous access control.

Hard Real-Time Support in seL4

- ▶ **Priority-Based Scheduling**
 - ▶ Simple, deterministic **priority-based policy**.
 - ▶ No autonomous priority adjustments—full control remains with the user.
- ▶ **Low Interrupt Latencies**
 - ▶ **Bounded latencies** achieved by
 - ▶ Disabling interrupts during kernel mode.
 - ▶ Simplified design with no need for complex concurrency control.
 - ▶ Efficient **short system calls** eliminate the need for preemptible kernels.

Handling Interrupt Latencies

There is a belief that a real-time OS must be preemptible in order to keep interrupt latencies low. This is true for traditional unprotected RTOSes running on simple microcontrollers. This is mistaken for a protected-mode system, such as seL4.

- ▶ When running on a powerful microprocessor with memory protection enabled, the time for entering the kernel, switching context, and exiting the kernel, is significant, and not much less than a seL4 system call
- ▶ In terms of interrupt latencies, little could be gained by a preemptible design, but the cost in terms of complexity would be very high, making a preemptible design unjustified
- ▶ This works as long as all system calls are short. In seL4 they generally are
- ▶ seL4 breaks operations into short sub-operations
- ▶ It's possible abort and restart the complete operation after each sub-operation, should there be a pending interrupt.

Worst-Case Execution Time (WCET) Analysis

▶ Key Features

- ▶ **Sound and complete WCET analysis** provides provable upper bounds for
 - ▶ System call latencies.
 - ▶ Interrupt handling.

Advantages for Real-Time Systems

1. Deterministic and efficient performance.
2. Responsive even during complex operations.
3. Guarantees critical for **safety-critical environments**.

Mixed-Criticality Systems (MCS)

Definition

A mixed-criticality system (MCS) is made up of (interacting) components of different criticalities.

Safety requirement: a failure of a component must not affect any more critical components, so the critical components can be assured independent of the less critical ones




Figure 7: A simplified example of mixed-criticality

Example

Avionics standards categorise failures from

- ▶ “no effect” (on vehicle safety) to
- ▶ “catastrophic” (loss of life).

The more critical a component, the more extensive (and expensive) is the required assurance, so there is a strong incentive for keeping criticalities low.

Traditionally,

critical systems would use a dedicated microcontroller for each function, i.e. isolation by air-gapping.

With growing functionality, this approach leads to a proliferation of processors (and their packaging and wiring), which causes **space, weight and power (SWaP)** problems, which MCS aim to overcome.

Key Characteristics of MCS

1. Strong Isolation

- ▶ Ensures lower-criticality failures cannot affect higher-criticality components.

2. Consolidation Goals

- ▶ Reduces **space, weight, and power (SWaP)** by consolidating functionality.
- ▶ Mimics the principle of isolating trusted and untrusted components with added safety requirements.

3. Real-Time Challenges

- ▶ Safety demands **timeliness** and real-time deadline adherence.
- ▶ Both **functional correctness** and timing are critical.

Traditional MCS with Time and Space Partitioning (TSP)

- ▶ **Strict Isolation:** guarantees for **spatial and temporal isolation**.
- ▶ **Resource Efficiency Issues**
 - ▶ **Worst-Case Execution Time (WCET)** allocations
 - ▶ Time slices sized for worst-case scenarios.
 - ▶ Leads to underutilized processor resources.
 - ▶ Slack time is wasted, as it cannot be reallocated.
- ▶ **Interrupt Latency Challenges**
 - ▶ Strict time slicing delays handling of **external events**.
 - ▶ Example Autonomous vehicle
 - ▶ Control loop runs every 5 ms.
 - ▶ Interrupts delayed by time slices, impacting responsiveness.

Trade-offs

1. **TSP** ensures strong isolation but resembles inefficiencies of air-gapped systems.
2. **Efficient MCS** requires balancing isolation with resource utilization and real-time responsiveness.

Mixed-Criticality Systems (MCS) in seL4

Core Challenge

Achieve **strong resource isolation** without the rigidity of strict Time and Space Partitioning (TSP).

Scheduling-Context Capabilities in seL4

▶ **Key Features**

- ▶ Regulate processor access by
 - ▶ **Time budget** How much CPU time a component can use.
 - ▶ **Time period** How often the budget can be used.
- ▶ Prevents components from monopolizing CPU while ensuring **responsiveness**.

▶ **Advantages over Traditional Time Slices**

- ▶ More **granular control** of CPU allocation.
- ▶ Enables dynamic resource utilization with strict isolation.

The MCS version of seL4

- ▶ replaces the time slice by a capability to a scheduling-context object
- ▶ Key to isolation: A scheduling context contains two main attributes.
 1. a *time budget*, which is similar to the old time slice, and limits the time for which a thread can execute until preempted.
 2. a *time period*, which determines how often the budget can be used: the thread will not get more time than one budget per period, preventing it from monopolising the CPU irrespective of its priority.

Example Critical vs. Non-Critical Components

▶ Scenario

- ▶ Critical control loop Requires **guaranteed CPU availability**.
- ▶ Non-critical driver Needs **high responsiveness** but must not interfere.

▶ Configuration

- ▶ Critical Controller
 - ▶ Budget **3 ms**.
 - ▶ Period **5 ms**.
 - ▶ Guarantees **60% CPU availability**.
- ▶ High-Priority Driver
 - ▶ Smaller budget **3 μ s**
 - ▶ Shorter period **10 μ s**
 - ▶ High responsiveness without exceeding **30% CPU time**.

▶ Result

- ▶ Ensures **critical deadlines** are met, regardless of non-critical behavior.
- ▶ Fulfills MCS requirements with flexibility and isolation.

Why seL4 for MCS?

- ▶ **Advanced Time Capabilities**

- ▶ Granular CPU control ensures **isolation** and **real-time guarantees**.
- ▶ State-of-the-art solution for **safety-critical environments**.

Deployment and Incremental Cyber Retrofit

Planning Deployment with seL4

1. Identify and Protect Critical Assets

- ▶ Structure assets as **modular, seL4-protected CAmkES components**.

2. Verification for Highest Assurance

- ▶ Use the verified kernel for your platform when possible.
- ▶ Even unverified versions provide stronger guarantees than most OSes.

3. User-Level Infrastructure

- ▶ Evaluate if existing components meet your needs.
- ▶ Collaborate with the community or specialized providers for missing infrastructure.

4. Contribute Back

- ▶ Share useful components under an **open-source license** to foster collaboration.

▶ Virtualization for Legacy Components

- ▶ Legacy systems often cannot be ported due to
 - ▶ Size or complex dependencies.
 - ▶ Minimal security benefits from running natively.
- ▶ Use seL4's **virtualization capabilities** as a baseline.

Example DARPA HACMS and Boeing ULB

- ▶ Initial Setup: Linux system placed in a **VM** on seL4.
- ▶ Incremental Transformation
 - ▶ Isolated untrusted components (e.g., camera software, GPS) into
 - ▶ Separate VMs.
 - ▶ Native CAMkES components
 - ▶ Critical modules moved to **secure native implementations**.




Figure 8: Incremental cyber-retrofit of the Boeing ULB mission computer

- ▶ Result: Even if Linux is compromised, the rest remains secure.

Key Benefits

- ▶ Gradual system modernization with **minimal disruption**.
- ▶ Enhanced **security and resilience** against attacks.
- ▶ Efficient reuse of legacy systems while isolating critical components.

Security is No Excuse for Poor Performance